

TeaStore

A Micro-Service Application for Benchmarking, Modeling and Resource Management Research

Jóakim von Kistowski, Simon Eismann, André Bauer, Norbert Schmitt,
Johannes Grohmann, Samuel Kounev




November 9, 2018

<https://github.com/DescartesResearch/TeaStore>



Example Research Scenario



Many solutions for these questions have been proposed, however...





How to evaluate


- Placement algorithms
 - Auto-scalers
 - New modeling formalisms
 - Model extractors
- Require realistic **reference and test applications**

Reference applications help to

- Evaluate model (extractor) accuracy
- Measure auto-scaler elasticity
- Measure placement power consumption and performance




Requirements for a Test Application

- Scalable
- Allows for changes at run-time
- Reproducible performance results
- Diverse performance behavior
- Dependable and stable
- Online monitoring 
- Load profiles
- Simple setup
- Modern, representative technology stack



Existing Test Applications


- RUBiS [1]
 - eBay-like bidding platform
 - Created 2002
 - Single service
- SPECjEnterprise 2010 [2]
 - SPEC Java Enterprise benchmark
 - Three tier architecture
 - No run-time scaling
 - Database is primary bottleneck
- Sock Shop [3]
 - Microservice network management demo application
 - Created 2016
 - Low load on non-network resources
- Dell DVDStore, ACME Air, Spring Cloud Demo, and more in our MASCOTS paper [4]




The TeaStore

Micro-service test application


- Five services + registry
- Netflix “Ribbon” client-side load balancer
- Kieker APM [5]
- Documented deployment options:
 - Manual
 - Docker images
 - Kubernetes



NETFLIX
OSS



kieker



Services I

Registry

- Simplified Netflix Eureka
- Service location repository
- Heartbeat



RegistryClient

- Dependency for every service
- Netflix “Ribbon”
- Load balances for each client



WebUI

- Servlets/Bootstrap
 - Integrates other services into UI
- **CPU + Memory + Network I/O**




Authentication

- Session + PW validation
 - SHA512 + BCrypt
- **CPU**



PersistenceProvider



- Encapsulates DB
- Caching + cache coherence

➤ **Memory**

Recommender



- Recommends products based on history
- 4 different algorithms

➤ **Memory or CPU**

ImageProvider



- Loads images from HDD
- 6 cache implementations

➤ **Memory + Disk I/O**

TraceRepository

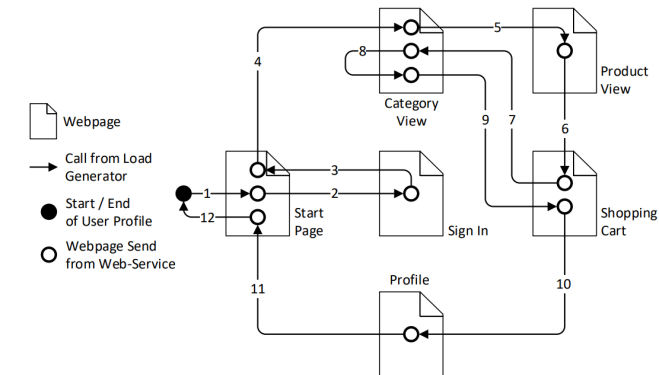



- AMQP Server
- Collects traces from all services



HTTP load generator [5]

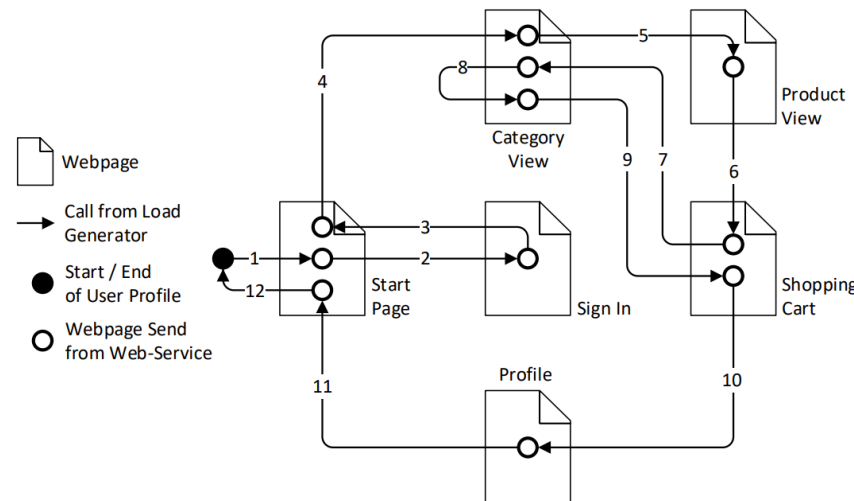

- Supports varying load intensity profiles
 - Can be created manually
 - Or using LIMBO [6]
- Scriptable user behavior
 - Uses LUA scripting language
 - “Browse” and “Buy” profiles on GitHub




<https://github.com/joakimkistowski/HTTP-Load-Generator>

JMeter

- Commonly used load generator
- Browse profile for JMeter
- Identical to HTTP Load Generator profile



Evaluation Teaser: Does it scale?






- Scales linearly
- Stresses 144 cores on 9 physical hosts
- HTTP Load Generator handles > 6000 requests per second



Three Use-Cases

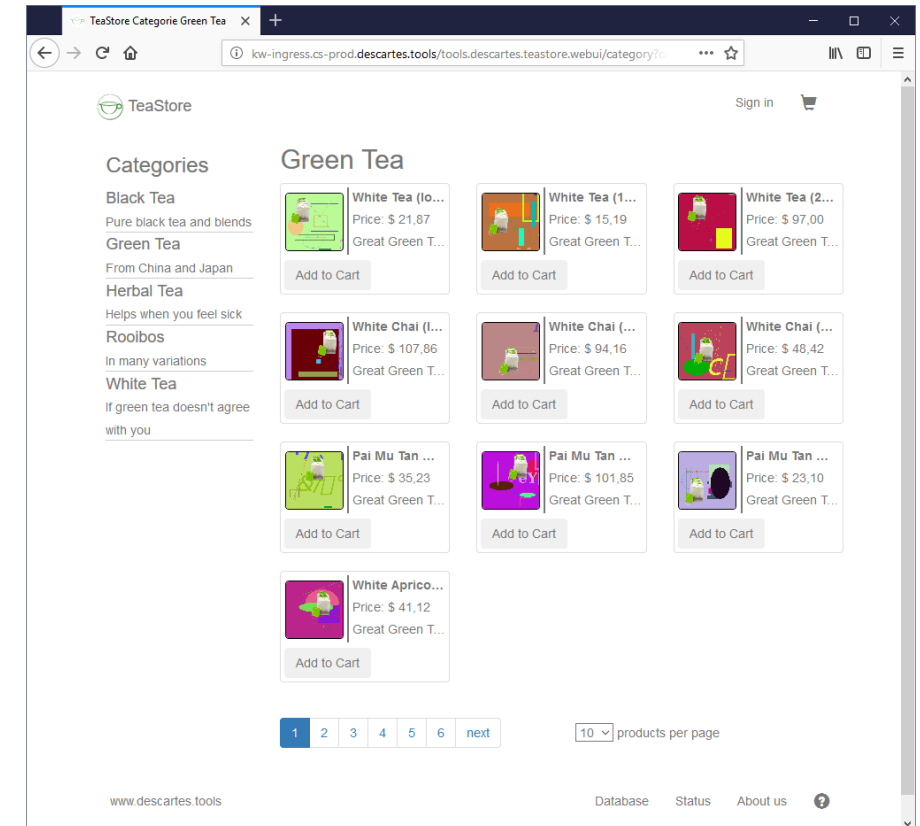
- Performance modeling
- Auto-scaling
- Measuring energy-efficiency of placements

Goal: Demonstrate TeaStore's use in these contexts



Performance Model - Scenario

- Question: How does utilization change with the default *# products per page* ?
- Approach:
 - Create two workloads with different *products per page* distributions
 - Create and calibrate performance model with default distribution
 - Predict performance for
 - Different *products per page* distribution
 - Different service placement



Products per Page Distribution

Calibration


$$P_5(x) = \begin{cases} 0.9 & \text{if } x = 5 \\ 0.09 & \text{if } x = 10 \\ 0.01 & \text{if } x = 20 \\ 0 & \text{else.} \end{cases}$$

To Predict


$$P_{10}(x) = \begin{cases} 0 & \text{if } x = 5 \\ 0.99 & \text{if } x = 10 \\ 0.01 & \text{if } x = 20 \\ 0 & \text{else.} \end{cases}$$

Deployment

Calibration




To Predict



Performance Model - Results


Results with and without considering the parametric dependency using Service Demand Law-based model



Auto-Scaling - Scenario

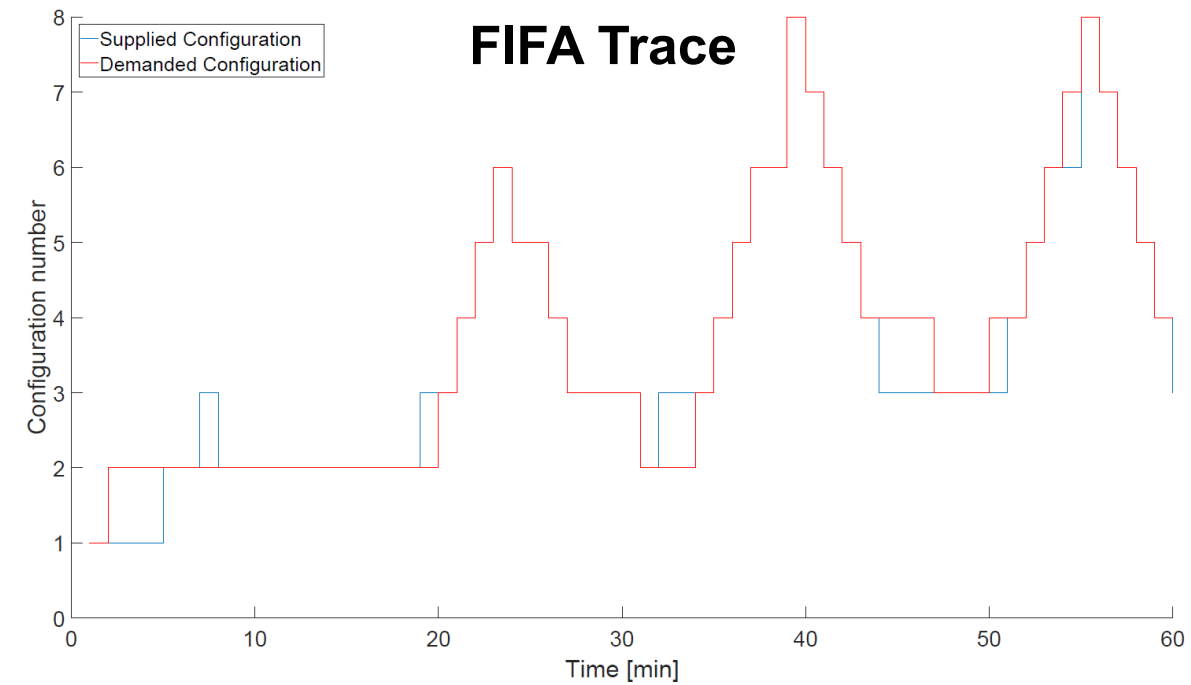
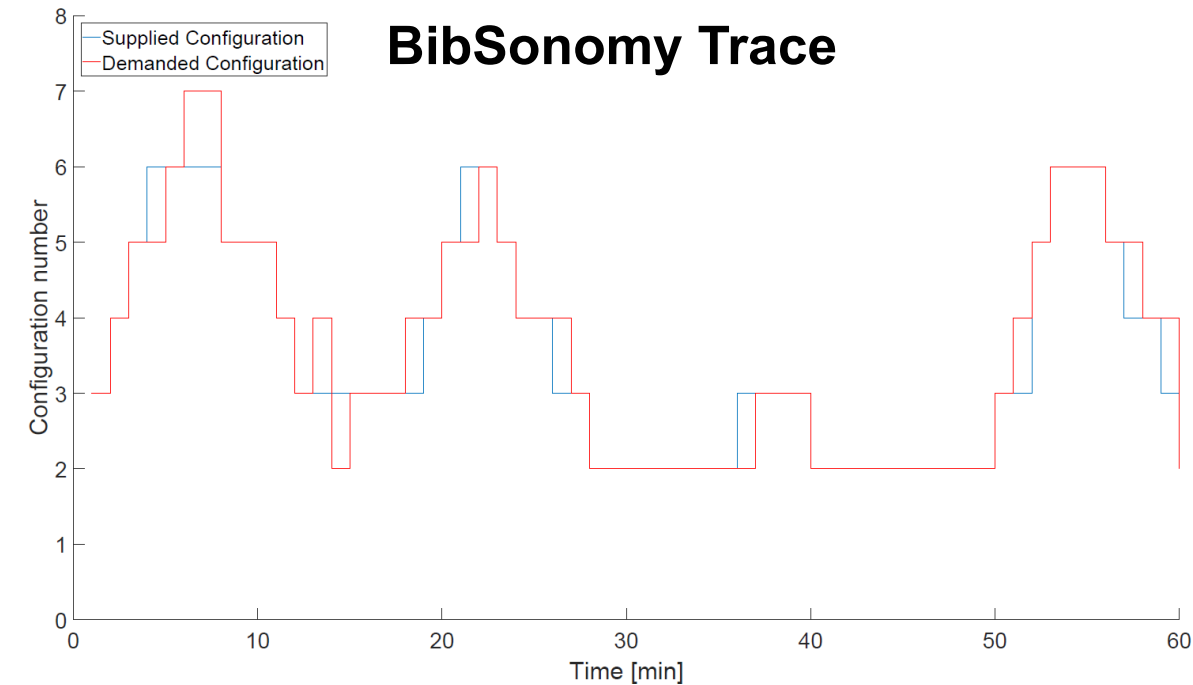
Reactive Auto-Scaling Scenario

- Challenge: Scale in an elastic manner so that *# services* matches demand
- Additional Challenge: Which service to scale?
- Approach:
 - Create heterogeneous configuration order
 - Put TeaStore under varying load
 - Decide scale-up / scale-down using research auto-scaler REACT [7]



Service	Configuration								
	#1	#2	#3	#4	#5	#6	#7	#8	#9
WebUI	1	1	2	3	3	4	5	5	6
Image Provider	1	1	2	3	3	4	5	5	6
Authentication	1	2	3	4	5	6	7	8	9
Recommender	1	1	1	2	2	2	3	3	3
Persistence	1	2	2	3	4	4	5	6	6

Auto-Scaling - Results



- Under- and overprovisioning-timeshare $\leq 15\%$
- TeaStore can be used for auto-scaler evaluation
- Open challenge: Which service to scale next?

Energy efficiency of placements


- Goal: Show that power consumption, energy efficiency, and performance scale differently
 - Different optima for service placements
- Approach:
 - Distribute TeaStore on homogeneous and heterogeneous servers
 - Put TeaStore under load using increasing stress-test load intensity
 - Measure TeaStore performance and server wall power



Energy Efficiency - Measurement

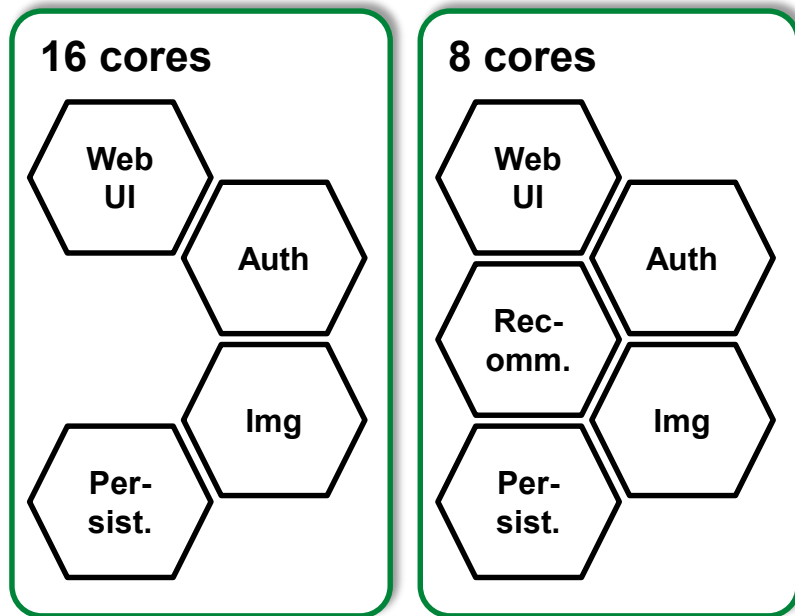
Measurements in homogeneous and heterogeneous setting

- SUT 1:
 - 16 core Haswell
 - 32 GB RAM
- SUT 2 (Heterogeneous):
 - 8 core Skylake
 - 16 GB RAM
- Metrics:
 - Throughput
 - Power
 - Energy Efficiency
 - Throughput / Power



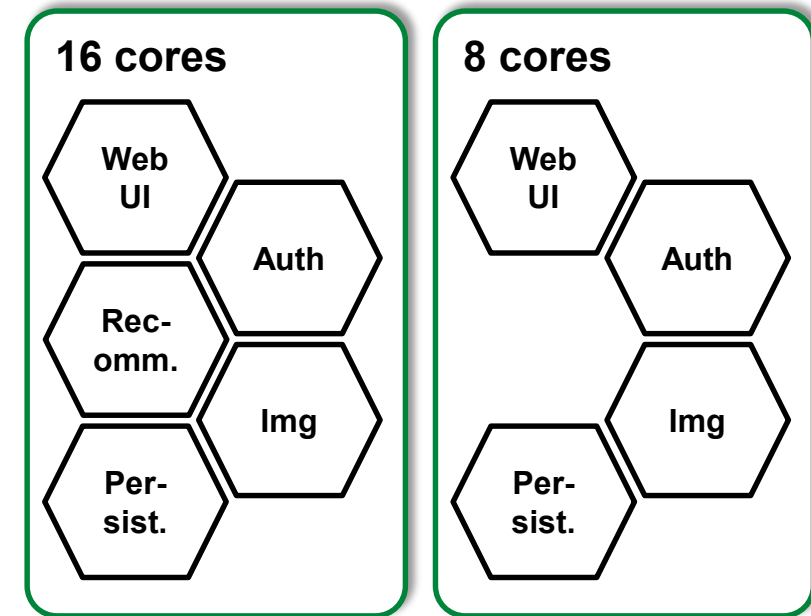
Energy Efficiency – Optima for Heterogeneous Placement

Placement Candidate 1



Max		1011.9 Tr/s
Max		179.6 W
Geo		4.4 Tr/J

Placement Candidate 2



Max		1067.7 Tr/s
Max		187.0 W
Geo		4.3 Tr/J

TeaStore - Conclusions

- Teastore can be used for
 - Performance modeling evaluation
 - Auto-Scaler evaluation
 - Placement and energy-efficiency evaluation
- Micro-service reference application
 - Five Services + Registry
 - Different resource usage characteristics
 - Kieker monitoring
 - Load Generators and Load Profiles
 - Kubernetes support
- Under Review by SPEC RG



<https://github.com/DescartesResearch/TeaStore>



Thank You!

From the TeaStore Dev Team



<https://github.com/DescartesResearch/TeaStore>

- HTTP Load Generator
<https://github.com/joakimkistowski/HTTP-Load-Generator>
- LIMBO Load Intensity Modeling Tool
<http://descartes.tools/limbo>
- Kieker Application Monitoring
<http://kieker-monitoring.net>



All tools available at:



<http://descartes.tools/>



References

- [1] RUBiS User's Manual, May 2008.
- [2] Standard Performance Evaluation Corporation (SPEC). SPEC jEnterprise 2010 Design Document. <https://www.spec.org/jEnterprise2010/docs/DesignDocumentation.html>, May 2010, Accessed: 16.10.2017.
- [3] Weaveworks Inc. Sock Shop: A Microservice Demo Application. <https://github.com/microservices-demo/microservices-demo>, 2017, Accessed: 19.10.2017.
- [4] A. van Hoorn, J. Waller, and W. Hasselbring. Kieker: A framework for application performance monitoring and dynamic software analysis. In *Proceedings of the 3rd joint ACM/SPEC International Conference on Performance Engineering (ICPE 2012)*, April 2012.
- [5] J. von Kistowski, M. Deffner, and S. Kounev. Run-time Prediction of Power Consumption for Component Deployments. In *Proceedings of the 15th IEEE International Conference on Autonomic Computing (ICAC 2018)*, Trento, Italy, September 2018.
- [6] J. von Kistowski, N. Herbst, S. Kounev, H. Groenda, C. Stier, and S. Lehrig. Modeling and Extracting Load Intensity Profiles. *ACM Transactions on Autonomous and Adaptive Systems (TAAS)*, 11(4):23:1 - 23:28, January 2017.
- [7] T. C. Chieu, A. Mohindra, A. A. Karve, and A. Segal. Dynamic scaling of web applications in a virtualized cloud computing environment. In *International Conference on E-Business Engineering*, 2009.
- [8] N. R. Herbst, S. Kounev, and R. Reussner. Elasticity in Cloud Computing: What it is, and What it is Not. In *Proceedings of the 10th International Conference on Autonomic Computing (ICAC 2013)*, San Jose, CA, June 2013.